

```
In [1]: import pandas as pd
import numpy as np
from ggplot import *
```

```
In [2]: # read input data - version 2
df2 = pd.read_csv("turnstile_weather_v2.csv", parse_dates=['datetime'])
# the original data
df = pd.read_csv('turnstile_data_master_with_weather.csv')
df2.head()
```

Out[2]:

	UNIT	DATEn	TIMEn	ENTRIESn	EXITSn	ENTRIESn_hourly	EXITSn_hourly	date1
0	R003	05-01-11	00:00:00	4388333	2911002	0	0	2011 01 00:00
1	R003	05-01-11	04:00:00	4388333	2911002	0	0	2011 01 04:00
2	R003	05-01-11	12:00:00	4388333	2911002	0	0	2011 01 12:00
3	R003	05-01-11	16:00:00	4388333	2911002	0	0	2011 01 16:00
4	R003	05-01-11	20:00:00	4388333	2911002	0	0	2011 01 20:00

5 rows × 27 columns


```
In [3]: df2.describe()
```

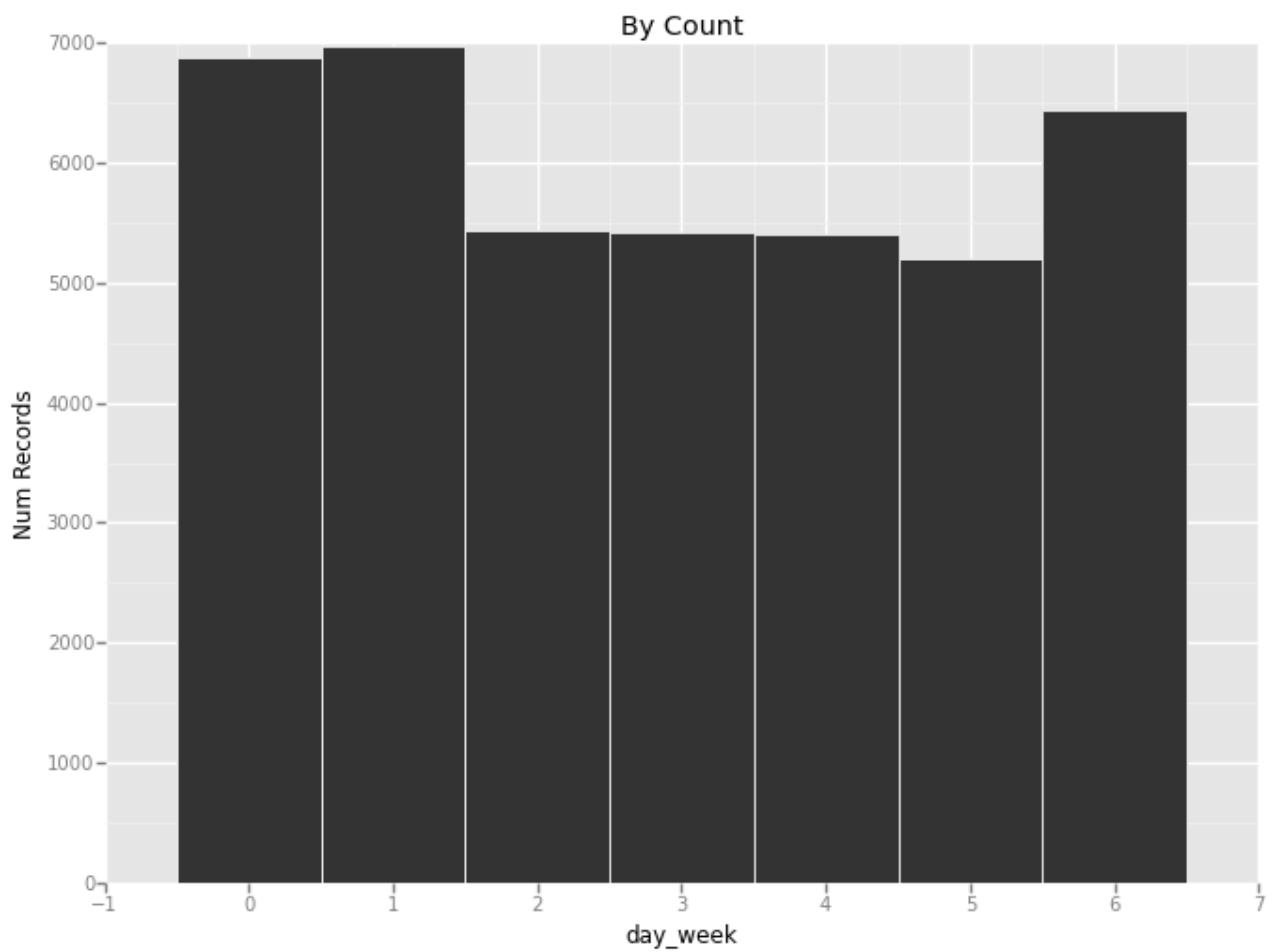
Out[3]:

	ENTRIESn	EXITSn	ENTRIESn_hourly	EXITSn_hourly	hour	d
count	4.264900e+04	4.264900e+04	42649.000000	42649.000000	42649.000000	4
mean	2.812486e+07	1.986993e+07	1886.589955	1361.487866	10.046754	2
std	3.043607e+07	2.028986e+07	2952.385585	2183.845409	6.938928	2
min	0.000000e+00	0.000000e+00	0.000000	0.000000	0.000000	0
25%	1.039762e+07	7.613712e+06	274.000000	237.000000	4.000000	1
50%	1.818389e+07	1.331609e+07	905.000000	664.000000	12.000000	3
75%	3.263049e+07	2.393771e+07	2255.000000	1537.000000	16.000000	5
max	2.357746e+08	1.493782e+08	32814.000000	34828.000000	20.000000	6

8 rows × 21 columns

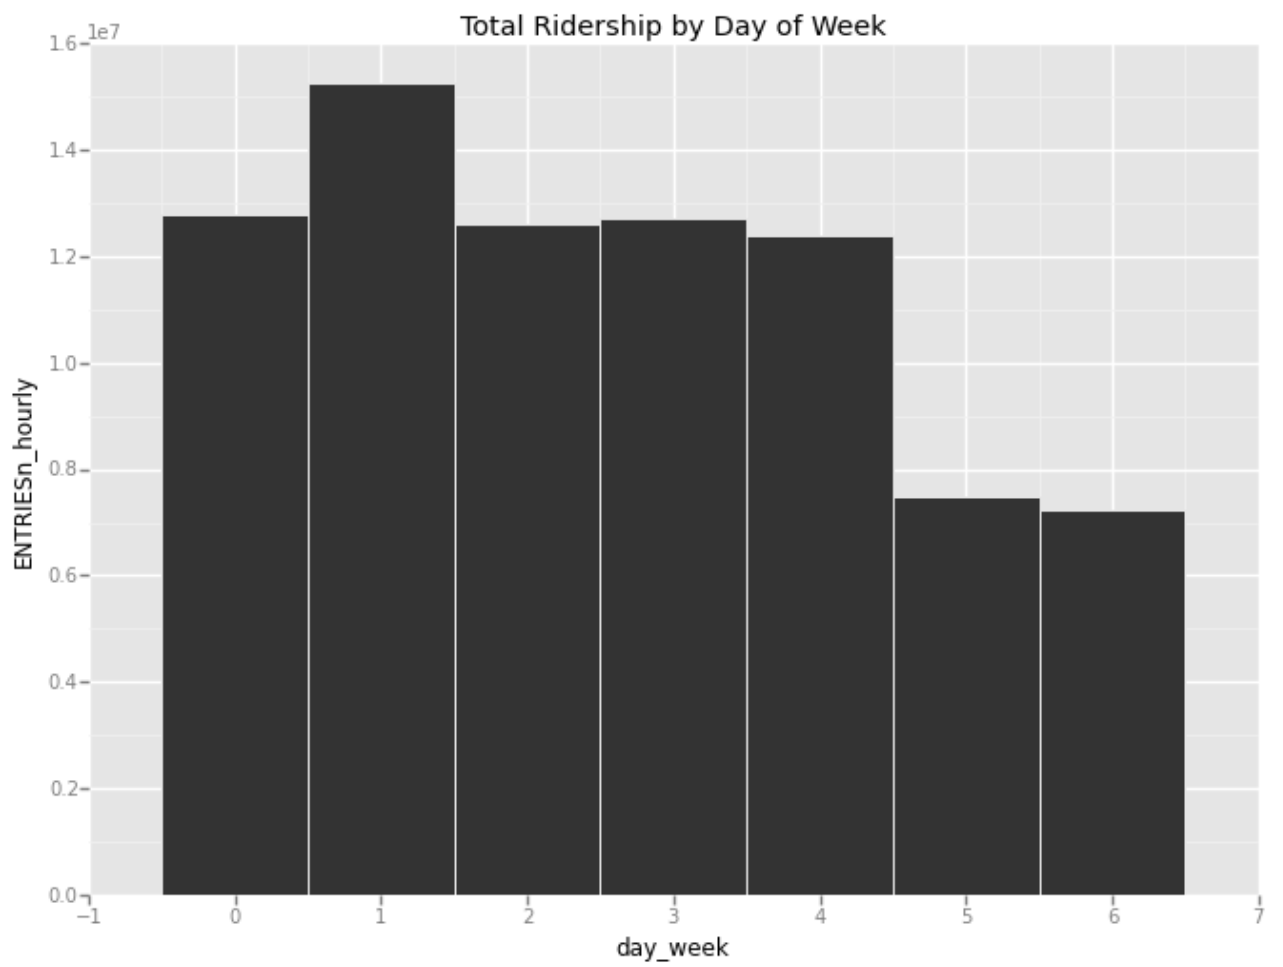
```
In [6]: # distribution of data by day of week
#ggplot(aes(x='day_week'), data=df2) + geom_histogram(binwidth=1) # note this is lumping last bin in with second to last
day_count = df2[['day_week', 'ENTRIESn_hourly']].groupby('day_week', as_index=False).aggregate(np.count_nonzero)
ggplot(aes(x='day_week', y='ENTRIESn_hourly'), data=day_count) + geom_bar(stat='identity') + labs(title='By Count') + ylab('Num Records')
```






Out[6]: <ggplot: (8738211923377)>

```
In [7]: day_sum = df2[['day_week', 'ENTRIESn_hourly']].groupby('day_week', as_index=False).aggregate(np.sum)
ggplot(aes(x='day_week', y='ENTRIESn_hourly'), data=day_sum) + geom_bar(stat='identity') + labs(title='Total Ridership by Day of Week')
```




Out[7]: <ggplot: (8738211773285)>

```
In [25]: # distribution of data by hour
hour_count = df2[['hour', 'ENTRIESn_hourly']].groupby('hour', as_index=False).aggregate(np.count_nonzero)
#ggplot(aes(x='hour'), data=df2) + geom_histogram(binwidth=1) # last bin is being included with next to last bin
ggplot(aes(x='hour', y='ENTRIESn_hourly'), data=hour_count) + geom_bar(stat='identity') + ylab('Num Records')
```



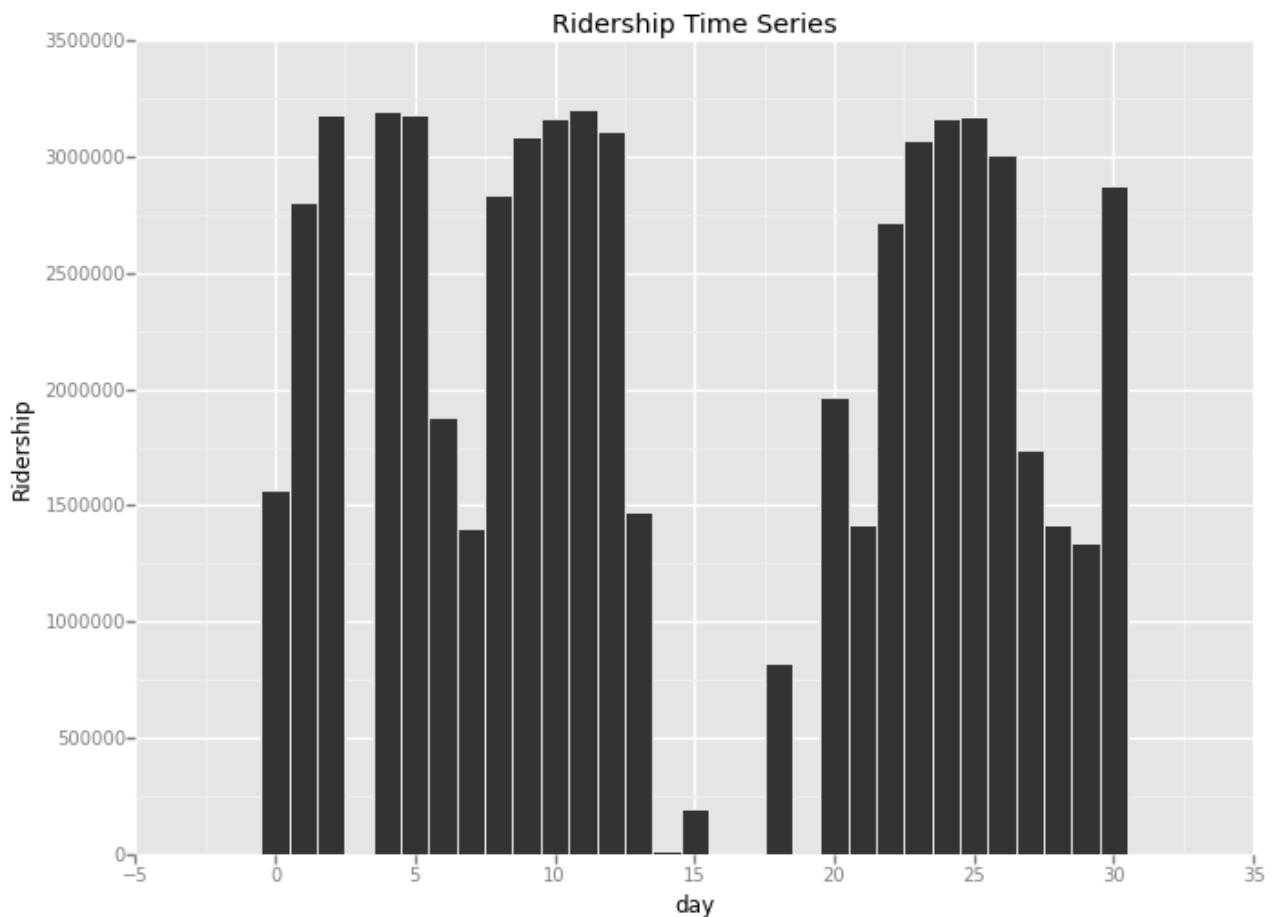
```
Out[25]: <ggplot: (8786094427405)>
```

```
In [8]: hour_sum = df2[['hour', 'ENTRIESn_hourly']].groupby('hour', as_index=False).aggregate(np.sum)
ggplot(aes(x='hour', y='ENTRIESn_hourly'), data=hour_sum) + geom_bar(stat='identity')
```



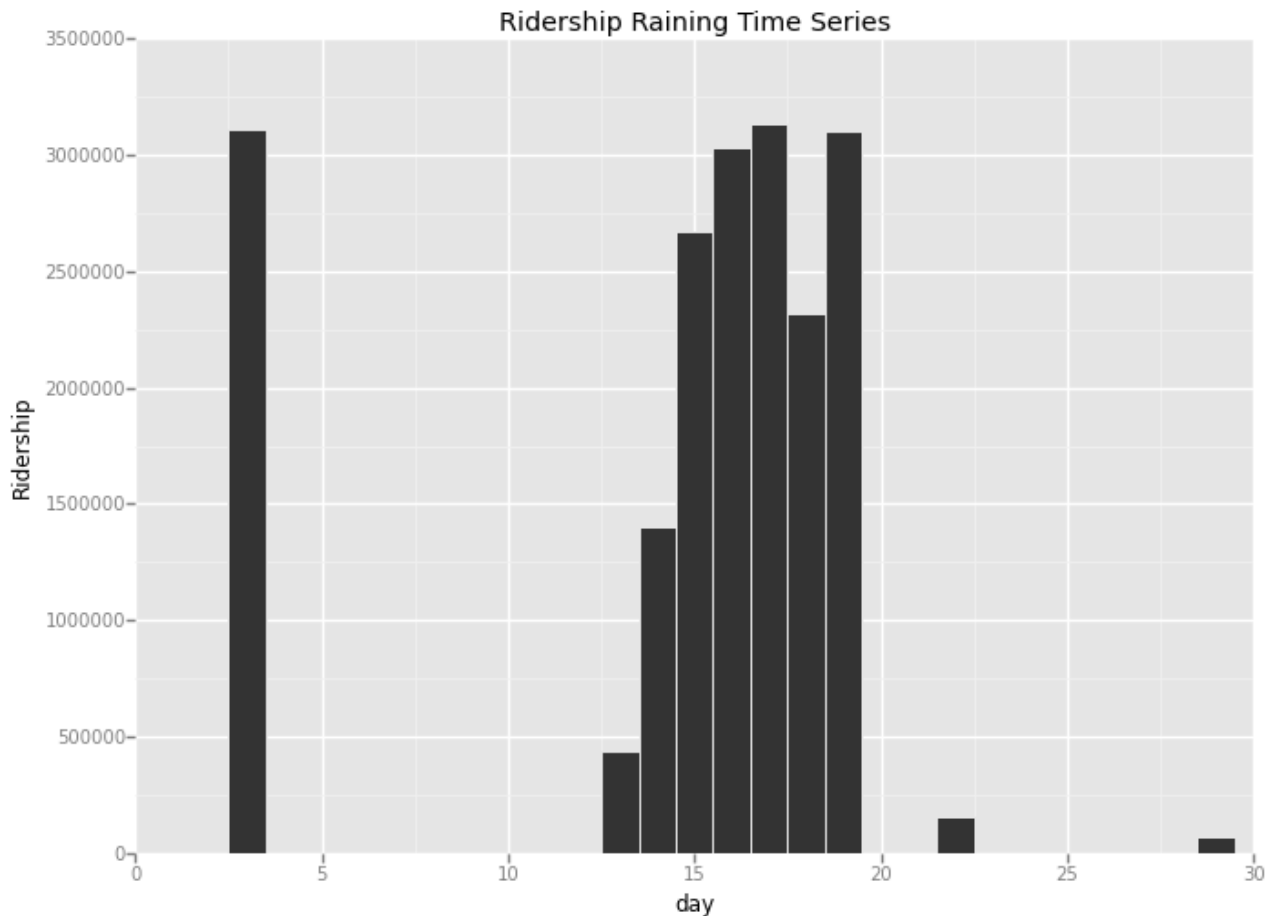
```
Out[8]: <ggplot: (8736946058553)>
```

```
In [19]: # let's plot as time series
# plotting the time series allows to spot gaps and outliers too
from datetime import datetime
norain = df2[df2['rain'] == 0]
epoch = datetime(2011, 5, 1)
# add seconds since some known time as column
norain['day'] = norain.apply(lambda x: (x['datetime'] - epoch).days, axis=1)
norain = norain.set_index(['day'], drop=False)
time_series = norain[['day', 'ENTRIESn_hourly']].groupby('day', as_index=False).aggregate(np.sum)
ggplot(aes(x='day', y='ENTRIESn_hourly'), data=time_series) + geom_bar(stat="identity") + ylab('Ridership') + labs(title='Ridership Time Series')
```



```
Out[19]: <ggplot: (8736945200025)>
```

```
In [27]: rain = df2[df2['rain'] == 1]
rain['day'] = rain.apply(lambda x: (x['datetime'] - epoch).days, axis=1)
rain = rain.set_index(['day'], drop=False)
time_series_rain = rain[['day', 'ENTRIESn_hourly']].groupby('day', as_index=False).aggregate(np.sum)
ggplot(aes(x='day', y='ENTRIESn_hourly'), data=time_series_rain) + geom_bar(stat="identity") + ylab('Ridership') + labs(title='Ridership Raining Time Series')
```



```
Out[27]: <ggplot: (8736944950769)>
```

```
In [28]: ggplot(aes(x='ENTRIESn_hourly'), data=df2) + geom_histogram(binwidth=500) + scale_x_continuous(limits=(0, 10000)) +\
facet_wrap('rain') + ylab('Num Records') + labs(title='Rain vs. No Rain Number of Records')
```



```
/home/jay/anaconda/lib/python2.7/site-packages/ggplot/ggplot.py:200: R
untimeWarning: Facetting is currently not supported with geom_bar. See
https://github.com/yhat/ggplot/issues/196 for more
```

```
information
```

```
warnings.warn(msg, RuntimeWarning)
```


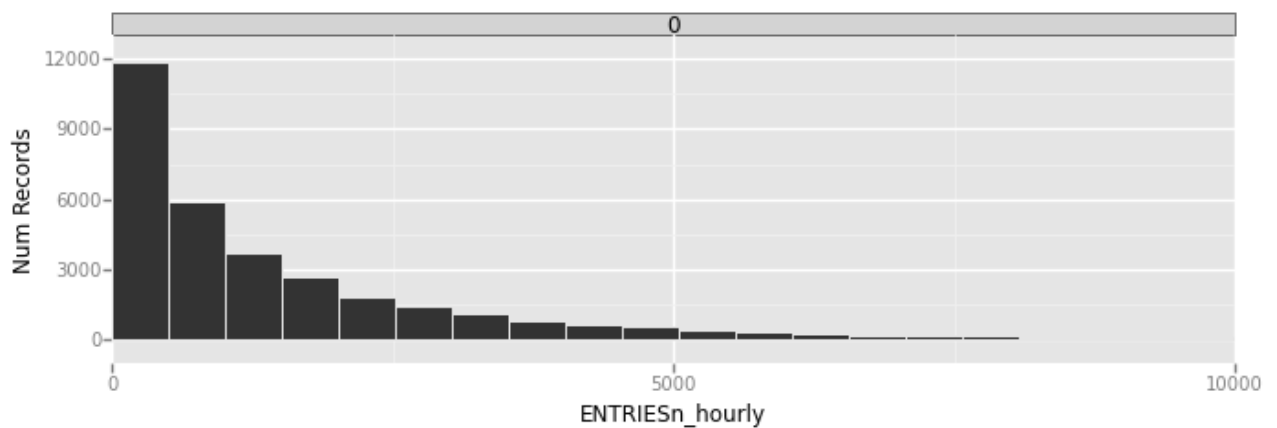
```
/home/jay/anaconda/lib/python2.7/site-packages/pandas/util/decorator
s.py:81: FutureWarning: the 'rows' keyword is deprecated, use 'index'
instead
```

```
warnings.warn(msg, FutureWarning)
```

```
/home/jay/anaconda/lib/python2.7/site-packages/ggplot/geoms/geom_bar.p
y:47: FutureWarning: comparison to `None` will result in an elementwis
e object comparison in the future.
```

```
_reset = self.bottom == None or (self.ax != None and self.ax != ax)
```

Rain vs. No Rain Number of Records



```
Out[28]: <ggplot: (8736944872637)>
```

```
In [10]: sum_rain = df2[['rain', 'ENTRIESn_hourly']].groupby('rain', as_index=False).aggregate(np.sum)
#ggplot(aes(x='ENTRIESn_hourly'), data=sum_rain) + geom_histogram(binwidth=500) + scale_x_continuous(limits=(0, 10000)) +\
#facet_wrap('rain') + ylab('Ridership') + labs(title='Rain vs. No Rain Ridership')
sum_rain
```

```
Out[10]:
```

	rain	ENTRIESn_hourly
0	0	61020916
1	1	19440259

```
In [41]: rain = df2[df2['rain'] == 1]['ENTRIESn_hourly']
rain.describe()
```

```
Out[41]: count      9585.000000
mean        2028.196035
std         3189.433373
min           0.000000
25%         295.000000
50%         939.000000
75%        2424.000000
max        32289.000000
Name: ENTRIESn_hourly, dtype: float64
```

```
In [42]: norain = df2[df2['rain'] == 0]['ENTRIESn_hourly']
norain.describe()
```

```
Out[42]: count      33064.000000
mean        1845.539439
std         2878.770848
min           0.000000
25%         269.000000
50%         893.000000
75%        2197.000000
max        32814.000000
Name: ENTRIESn_hourly, dtype: float64
```

```
In [43]: print rain.mean() - norain.mean()
print rain.median() - norain.median()

182.656596808
46.0
```

```
In []:
```